

# RL211 - HW4

Omri Himelbrand - 200863843

Nitzan Cohen - 203980750

January 7, 2021

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Running the solution</b>	<b>1</b>
2.1	Running as a script . . . . .	1
2.2	Running as module . . . . .	2
<b>3</b>	<b>Code details</b>	<b>2</b>
3.1	Global variables . . . . .	2
3.2	The functions . . . . .	2
<b>4</b>	<b>Plot</b>	<b>4</b>

## 1 Introduction

The code uses Actor-Critic algorithm to approximate the value function  $\hat{v}(s, W)$  and the policy  $\pi(a|s, \theta)$ . Using LFA for approximating  $v(s)$  with our set of weights  $W$  along with extracted feature-vector representing the state and  $\pi(a|s)$  with our set of weights  $\theta$  for each action given the state along with soft-max. Denote  $\alpha^W, \alpha^\theta$  as learning rates for updating  $W, \theta$  respectively, the code implements Auto-Step-Size Selection method for  $\alpha^W$  and  $\alpha^\theta$  separately, which theoretically proposed by us and empirically tested to work and converge. Denote  $v_{best}(0)$  as the highest  $v(0)$  value so far through the algorithm run, after each run of Actor-Critic algorithm and evaluation of  $v(0)$ , we calculate the  $ratio = v(0)/v_{best}(0)$  and multiply  $\alpha^W, \alpha^\theta$  using  $ratio^2, ratio^3$ , respectively, upper bounded by the initial  $\alpha^W, \alpha^\theta$ , though implementing Auto-Step-Size Selection.

## 2 Running the solution

### 2.1 Running as a script

```
usage: hw4.py [-h] [-human] [-gamma G] [-d] [-ms MAX_STEPS] [-es EVAL_STEPS]
              [-png PNG_SUFFIX] [-relax]
```

Actor - Critic **for** AI-Gym Mountain Car.

optional arguments:

```
-h, --help          show this help message and exit
-human             use this flag to run human agent
-gamma G          a float gamma in [0,1] (default: 1.0).
-d                use this flag to get debug prints
-ms MAX_STEPS     a int for number of maximum steps for learning.
-es EVAL_STEPS    a int for number of steps between evaluations.
-png PNG_SUFFIX   a suffix for png out file
-relax            use this flag to use 500 steps episodes
```

## 2.2 Running as module

```
import hw4

hw4.main()
```

## 3 Code details

### 3.1 Global variables

**DEBUG:** A boolean which is initialized to False.

**MAX\_STEPS:** An integer to indicate the number of steps for the entire learning process, initialized to 150000.

**EVAL\_STEPS:** An integer to indicate the number of steps between running evaluations simulations, initialized to 500.

**RELAXED:** A boolean value, indicating whether or not to run the relaxed version (500 steps limit per episode).

**P\_CENTERS:** A list containing the centers of the positions' Gaussians.

**V\_CENTERS:** A list containing the centers of the velocities' Gaussians.

**CENTER\_PRODUCTS:**  $P\_CENTERS \times V\_CENTERS$ .

**SIGMA\_P:** variance/std of position.

**SIGMA\_V:** variance/std of velocity.

**COV:** the co-variance matrix.

**INV\_COV:**  $COV^{-1}$

**P\_I:** float value of interval size for centers of position.

**V\_I:** float value of interval size for centers of velocity.

*#initializing the last few globals can be done using these functions*

```
def init_intervals(Ip=0.18,Iv=0.014):
    global P_I,V_I
    P_I = Ip
    V_I = Iv

def init_covariance(sigma_p=0.04,sigma_v=0.0004):
    global SIGMA_P,SIGMA_V,COV,INV_COV
    SIGMA_P = sigma_p
    SIGMA_V = sigma_v
    COV = np.diag([SIGMA_P,SIGMA_V])
    INV_COV = np.linalg.inv(COV)

def init_centers(p_half=4,v_half=4):
    global P_CENTERS,V_CENTERS,CENTER_PRODUCTS
    P_CENTERS = [(i + 0.5)*P_I for i in range(-p_half,p_half)]
    V_CENTERS = [(i)*V_I for i in range(-v_half,v_half)]
    CENTER_PRODUCTS = np.array(list(product(P_CENTERS,V_CENTERS)))
```

### 3.2 The functions

```
def init_env(max_steps=200):
```

Calls gym.make, sets the max episode steps for the created environment and returns it.

```
def set_debug(value):
```

Sets the global variable DEBUG to value.

```
def set_max_steps(value):
```

Sets the global variable MAX\_STEPS to value.

```
def human_agent():
```

Prompts user to pick action for the next step of simulation, used for mostly for debugging.  
Return value is a action (int). Note that in order to use this you must have "readchar" installed.

```
def evaluate(env,w,gamma,episodes_num=100,show=False):
```

Given the current weights ( $w$ ), and the rest of the arguments seen in the signature, evaluates  $v_0^\pi$ , using a MC-like evaluation.

This function returns value  $v_0^\pi$ .

```
def apply_policy(Qhat,actions,eps=0):
```

This function apply the policy defined by the weights using  $\epsilon$ -greedy scheme on  $\hat{Q}$  (in case  $eps=0$ ), generating random number, if it is less than  $\epsilon$  uniform random choice of action from action space, else argmax.

```
def AC(env, w, theta, gamma, alphas, actions, eps, max_step=5000, iters=0,\
        epsilon_decay=0.999, min_eps=0.1):
```

This function does Actor-Critic with LFA, as described in the introduction section.

```
def run_simulation(env,theta=None,human=False,show=True):
```

resets env to initial state, then runs simulation either using the given policy (by weights) or using a human agent.

```
def centers_distance(s:np.ndarray):
```

given a state returns vector of differences from the CENTER\_PRODUCTS vector.

```
def state_features(x:np.ndarray):
```

calls *centers\_distance*(x) and uses the returned value to computes the features of the state.

```
def init_weights(nA=3,seed=27021990):
```

Initializes the weights randomly using a seed.

```
def piApproximation(theta:np.ndarray,s:np.ndarray):
```

given  $\theta$  and  $s$ , computes  $\pi(\cdot|s, \theta)$ .

```
def VApproximation(s: np.ndarray, w: np.ndarray):
```

given  $w$  and  $s$ , computes  $\hat{V}(s, W)$ .

```
def learn_policy(env, actions, gamma):
```

This function runs the whole learning process, running AC for EVAL\_STEPS, then running evaluation.

After each evaluation the  $V_{init}^\pi$  is saved with the number of total steps taken so far, and we keep track of the best  $\theta$  according to the evaluations so far, keeping it for return.

The return value of this function is: x - array of step counts, y - array of  $V_{init}^\pi$  collected, and the best  $\theta$ .

```
def main(gamma=1,human=False):
```

This function is called when running the code as a script, but can be used as seen above, this function does the following:

- calls "init\_env", "init\_covariance", "init\_intervals", and "init\_centers"
- if human flag is set, runs a single simulation with a human agent.
- else calls "learn\_policy", after running a simulation using a previous learned weights (if such exists).
- calls "run\_simulation" using the returned  $\theta$
- After running the learning process and collecting all the x and y values, calls plot\_results function

## 4 Plot




Figure 1: Original problem 200 steps limit per episode